



CMUX x AIM Hackathon

SUNDAY, 26TH APRIL

8 AM — 8 PM · Seoul, South Korea

AttentionX ▲ Vercel ANTHROPIC



Document outline

Document outline.....	2
Schedule.....	4
About the sponsors.....	4
Prizes & Categories.....	6
Track Judges.....	7
Sejun Kim.....	7
Jun Kim.....	7
Kyungseok Hur (허경석).....	7
Sanguen Chang.....	7
Suho Park.....	7
Kyungsoo Kim.....	8
Juhyun Song.....	8
Dasol Choi.....	8
Yeeun Kim.....	9
Pia Park.....	9
Dan Han (한관엽).....	9
Jiwoo (Chloe) Lee님.....	9
Taeung Kang.....	10
Seonmin Lee.....	10
Vadim Choi.....	10
Final Judges.....	11
Haon Park.....	11
Austin Wang.....	11
Sukone Hong.....	11
Event miscellaneous.....	12
Location	
모나코스페이스.....	12
Credits.....	12
Vercel.....	12
Claude.....	12
Gemini.....	12
Venue Wifi.....	12
Tracks.....	13
🔧 Developer Tooling.....	13
💡 Business & Applications.....	13
🔴 AI Safety & Security.....	13
Judging Criteria.....	14
🔧 Developer Tooling.....	14
💡 Business & Applications.....	14
🔴 AI Safety & Security.....	15
Submission.....	16
Submission form.....	16

Judging format.....	16
Rules.....	17

Schedule

08:00 Doors Open · Breakfast

08:30 Opening

09:00 Hacking Starts

13:00 Lunch

18:00 Submissions Close · Dinner · Round 1 Judging Begins

19:25 6 Finalists Announced

19:30 Final Pitches

20:00 Winners Announced

About the sponsors

About [Manaflow](#) (Makers of [cmux](#))

Manaflow is an open source applied AI lab building next-generation developer tools for the agentic era. Their flagship product, cmux, is an open source terminal built specifically for coding agents - giving AI the native execution environment it needs to write, run, and iterate on code effectively.

About [AIM Intelligence](#)

AIM Intelligence is an enterprise AI security company that helps organizations keep their AI systems safe and under control. Their two flagship products - Stinger, an automated red teaming tool, and Starfort, a real-time guardrail system - protect against AI vulnerabilities across text, image, audio, and video. Trusted by the likes of OpenAI, Microsoft, Anthropic, and major Korean banks, AIM Intelligence backs its products with published research from top AI conferences and recently raised a \$7M Series A.

About [AI Nexus](#)

AI Nexus is an event agency with a global community of 5,000+ builders passionate about AI. They run hackathons and tech events across 6 locations worldwide, currently, with the leading AI innovators.

About [AttentionX](#)

AttentionX is an AI research and startup group based in Seoul, bringing together builders from academia and industry across top institutions including Berkeley, Stanford, KAIST, Google DeepMind, and more. They work on cutting-edge research, open-source projects, and real-world applications, with published work at NeurIPS, ICML, ACL, and EMNLP - and a portfolio of startups backed by Y Combinator, a16z, and leading Asian venture funds.

Prizes & Categories

Grand Prize:

- 3,000,000원

Developer Tooling:

- 1,000,000원

Business & Applications:

- 1,000,000원

AI Safety & Security

- 1,000,000원

Student Special

- \$5000 Azure credits

Track Judges

Sejun Kim

Sejun Kim is the CEO of E Corp. and the operator of socialseed.ing, an AI-powered influencer marketing platform. He focuses on building AI developer tools, automation workflows, and agent-first products, primarily using Go, TypeScript, Next.js, and Python. In 2026, he was selected for Kakao MCP Player 10 and won 1st Place at the DigitalOcean Gradient AI Hackathon.

<https://www.linkedin.com/in/combba/>

Jun Kim

Jun Kim is the co-founder and CEO of Aside, an AI-powered browser that works on your behalf, backed by Y Combinator (F25). Before Aside, he was a founding member of Airbridge (AB180), Korea's leading MarTech company, where he joined at 17 as the first team member.

<https://www.linkedin.com/in/therne>

Kyungseok Hur (허 경 석)

CEO of PAYAP, a commerce platform built for small online sellers in Vietnam, with a long-term vision for Southeast Asia. Focused on helping merchants easily launch and operate their own online stores in a fast-changing digital market. Previously led business development for Kakao Mobility's taxi and designated driver businesses, with hands-on experience in operations and partnerships.

<https://www.linkedin.com/in/kyungseok-hur-b41629117>

Sanguen Chang

Sanguen Chang is the Co-Founder of E Corp., a serial entrepreneur who has built 9+ products across AI, software, and real-world applications. He focuses on AI agents, developer tooling, and practical AI systems, with experience spanning Python, Next.js, React, LangChain, RAG, and TensorFlow. He has built and launched multiple award-winning hackathon projects and enjoys turning technical ideas into usable products.

<https://www.linkedin.com/in/sgwannabe/>

Suho Park

Suho Park is a Product Manager at AIM Intelligence where he focuses on responsible AI deployment for enterprise. Previously at Shinhan Investment Corp, he led the AI Private Banker (AI PB) — a conversational AI and realtime financial content providing service for retail investors on the Shinhan SOL Securities platform - scaling it to 290K+ monthly active users with 54% monthly cohort retention. He also led Alli, a B2B AI agent platform serving 24 enterprise clients in Japan including Mizuho Financial Group.

<https://www.linkedin.com/in/suho-park-637908194>

Kyungsoo Kim

I'm Kyungsoo Kim, a Staff AI Product Designer at KT, working in the Agentic AI Lab. I focus on designing AI-powered services and conversational agents, bridging user experience design with large language models and emerging AI technologies. Previously, I worked as a Senior Product Design Lead at NCSoft, where I led AI-driven product initiatives including PAIGE, an AI-based baseball service that received an iF Design Award.

I'm particularly interested in practical AI applications, user-centric design, and building scalable AI agent experiences.

<https://www.linkedin.com/in/kimkyungsoo>

Juhyun Song

Juhyun Song is a security researcher at the KAIST Hacking Lab specializing in automated security assessment of system software, including TEEs, operating systems, and DBMSs. His work has appeared in ACM CCS and IEEE Access. Previously, he conducted TrustZone security research at Samsung Electronics and was named to the Top 10 (Hall of Fame) of KITRI's "Best of the Best," a flagship national program in Korea for developing elite cybersecurity talent.

<https://www.linkedin.com/in/juhyun-song-267231395/>

Dasol Choi

Principal Researcher, AIM Intelligence

Dasol Choi is an AI safety evaluation and red-teaming expert and a Principal Researcher at AIM Intelligence. With a rigorous research background spanning LLMs, VLMs, and audio models, Dasol has published significant works in top-tier venues on evaluating AI systems across diverse modalities, focusing on alignment, bias mitigation, and defending against complex adversarial attacks. As the Lead Organizer of the 'Judgement Day' AI red-teaming arena, Dasol bridges the gap between theoretical AI safety research and real-world deployment.

<https://www.linkedin.com/in/dasol-choi-570797261/>

Yeeun Kim

Yeeun Kim is an AI Engineer at KT, working on the development and evaluation of LLM-based agents within MyK Agent, KT's commercial AI service launched in 2025. She focuses on building reliable and scalable agent systems through RAG, multi-turn optimization, and evaluation frameworks.

Her M.S. research in Artificial Intelligence focused on LLM safety and benchmark design, with publications at AAAI 2025, IJCAI 2024 Workshop, and EMNLP 2024 (KBL). Her recent work extends to RLVR-based safety alignment and robustness.

<https://www.linkedin.com/in/yeahrl0>

Pia Park

<https://www.piapark.me/>

Pia Park is an independent systems and cryptography engineer recent work includes secure LLM sandbox, security awareness, and infra. She primarily worked in Rust, also with Python, Typescript. Background in cryptographic engineering across zkp, iO, FHE, including projects funded by Ethereum Foundation and ARIA Research. She has won several global scale hackathons, including 1st place at the Europe Defense Tech Hackathon London.

<https://www.linkedin.com/in/pia-park/>

Dan Han (한관엽)

CEO, Co-founder @ Aleph Lab (YC F25)

<https://www.linkedin.com/in/gy-han/>

Jiwoo (Chloe) Lee

I am a Ph.D. Candidate in Materials Science & Engineering at Yonsei University and Imperial College London, currently specializing in machine learning-driven dopant selection in semiconductors.

Beyond my research, I am the founder of 가방끈 (pronounced "Gabang-ggeun.", meaning "bag strap" in English), a South Korean community application for graduate students, currently with over 50,000 Total Users. 가방끈 is dedicated to fostering the professional development of researchers and facilitating job market connections for master's and doctoral candidates.

<https://www.linkedin.com/in/bagstrap-jiwoo>

Taeung Kang

Taeung Kang is a frontend engineer specializing in user experience, performance optimization, and system architecture across healthcare, platform, and community domains. He has led zero-to-one product development at multiple startups, including B2B healthcare SaaS and cloud EMR systems. As a co-founder of Driving Teacher, he built and scaled an MVP to nationwide adoption and secured significant seed funding. Taeung combines deep technical expertise in React, TypeScript, and frontend architecture with hands-on experience in design systems, CI/CD, and product strategy. With entrepreneurial and leadership experience, he focuses on building impactful, scalable products that deliver meaningful value to users.

<https://www.linkedin.com/in/강태웅>

Seonmin Lee

Seonmin Lee (@LilMGenius) is the CEO/CAIO of defytheodd (d8d), an applied AI systems integration company running 30+ enterprise AX projects across 15+ verticals, backed by top VCs. Self-taught from rural Gangwon-do since age 9, he shipped a 3,000-player MMORPG at 12 (before learning relational databases existed), won silver at the Korean Olympiad in Informatics at 14, then spent ages 15 to 20 as a hip-hop artist. The youngest-ever AI Principal at a major Korean fashion conglomerate (~\$7B revenue), where he solo-shipped 10+ production AI systems in 6 months. Won Ralphthon Seoul #2 (Sponsored by OpenAI) as a 1-person team with 0 keyboard touches after the initial Enter, while other finalists logged 3-4+. The harness did the work. That harness, Polysona, is now open-source.

<https://www.linkedin.com/in/lilmgenius/>

Vadim Choi

CTO @ Leviosa AI

Selected for Antler Korea Cohort 8, chose to build full-time instead.

Hackathon winner (1st @ GSC, 2nd @ AngelHack).

<https://www.linkedin.com/in/vadimtsoi>

Final Judges

Haon Park

CEO, AIM Intelligence

<https://www.linkedin.com/in/haonpark>

Austin Wang

Founder & CEO of Manaflow

cmux.com

manaflow.com

austinywang.com

<https://www.linkedin.com/in/austinwang115>

Sukone Hong

Co-President & Fund General Partner of AttentionX. We are an AI research & startup group based in Seoul from academia and industry. We work on research, open-source projects, and real-world applications to build the next leading AI startups. We also invest in startups. Read more about our venture fund [here](#).

<https://attentionx.org/>

<https://www.linkedin.com/in/my-name-is-sukone>

Event miscellaneous

Location

📍 모나코스페이스

Credits

Vercel

\$30 v0 credits - use code: V0-MANAFLOW - <https://v0.app/>

Claude

\$25 - https://claude.com/offers?offer_code=5bd9a20b-2b47-47ba-9e64-b82db0cc01cf

Gemini

Gemini api key #1:

AIzaSyDk_jf7hN5z7Pdi4ohLgInK8sMgkgpph2c

Gemini api key #2:

AIzaSyCYcJpi4o1bbtXa1S3HQBnJN1m7iZt2DhA

Gemini api key #3:

AIzaSyBB2tGU2DcpmpbRfWsan00C37_deP13vug

Gemini api key #4:

AIzaSyBMPg2CYgB7MK4QpMeQLKZ-kNrR-xgvY-o

Venue Wifi

1)

Name: monacospace_main_5G

Password: main1234

2)

Name: Hack
Password: 1234567!

Tracks



Developer Tooling

CLI-first developer tools. Terminal workflows, agent orchestration, code generation, debugging.



Business & Applications

Build and deploy an AI-powered product in one day.



AI Safety & Security

Find, expose, and defend against AI vulnerabilities. Red teaming CLIs, guardrail pipelines, jailbreak detection, prompt injection scanners, LLM fuzzing. AIM Intelligence engineers on-site to mentor and judge.

Judging Criteria

Developer Tooling

CLI-first developer tools. Terminal workflows, agent orchestration, code generation, debugging.

Criterion	1 - Surface	2 - Functional	3 - Solid	4 - Impressive	5 - Exceptional
Technical Depth 30%	Basic wrapper around existing tools, no novel engineering	Working integration with some custom logic	Non-trivial engineering, handles edge cases	Complex architecture, custom protocols or parsers	Novel technical contribution, publication-worthy depth
Developer Experience 25%	Hard to install, confusing UX, poor feedback	Works but has friction; docs are incomplete	Smooth CLI/API, reasonable defaults, helpful errors	Thoughtful ergonomics, great error messages, fast	"This is how it should work." Immediately intuitive
Usefulness & Real-World Fit 25%	Unclear who benefits or when you'd use it	Useful in narrow, specific scenarios only	Solves a real daily pain point for developers	Many devs would reach for this immediately	Fills an obvious gap; hard to unsee it
Demo & Presentation 10%	Hard to follow what the tool does	Shows it works but misses the compelling use case	Concrete demo, explains the value well	Live demo lands; judges want to try it themselves	Story + demo make the case unforgettably
Judge's Personal Rating 10%	Not interesting or compelling to me personally	Somewhat interesting but wouldn't think about it again	Solid idea; I can see the appeal	I'd tell a colleague about this	This genuinely excites me; I want to see it succeed

Business & Applications

Build and deploy an AI-powered product in one day.

Criterion	1 - Vague	2 - Generic	3 - Defined	4 - Sharp	5 - Compelling
Problem & Market Clarity 25%	No clear problem statement; solution looking for a problem	Problem exists but is broad; no specific user pain	Clear problem, identifiable target users	Specific pain point with evidence; defined ICP	Urgent, large, underserved problem with strong insight
Product Completeness 20%	Slide deck or mockup; nothing functional	Works in one happy path; many rough edges	Core loop works end-to-end; usable in a real scenario	Handles multiple flows, error states, and real data	Could be handed to a user today. Impressive for 24h
AI Integration 20%	AI is a gimmick; product works without it	AI enhances one feature but isn't core	AI is central to the product experience	Product couldn't exist without AI; unlocks new behavior	AI changes the paradigm; creates a 10x improvement
Business Viability 15%	No thought given to monetization or sustainability	General idea ("charge SaaS fees") but no specifics	Clear model with realistic unit economics sketched	GTM + pricing + competitive moat addressed	Makes a compelling case for traction and growth
Demo & Presentation 10%	Hard to follow what the product does	Shows it works but misses the compelling use case	Concrete demo, explains the value well	Live demo lands; judges want to try it themselves	Story + demo make the case unforgettably
Judge's Personal Rating 10%	Not interesting or compelling to me personally	Somewhat interesting but wouldn't think about it again	Solid idea; I can see the appeal	I'd tell a colleague about this	This genuinely excites me; I want to see it succeed

AI Safety & Security

Find, expose, and defend against AI vulnerabilities. Red teaming CLIs, guardrail pipelines, jailbreak detection, prompt injection scanners, LLM fuzzing. AIM Intelligence engineers on-site to mentor and judge.

Criterion	1 - None	2 - Vague	3 - Defined	4 - Rigorous	5 - Expert
Threat Model & Rigor 30%	No clear attack or vulnerability being targeted	Identifies a class of risk but no specific threat model	Clear threat scenario with attacker goals and assumptions	Formal or structured threat model; edge cases considered	Publication-quality threat analysis; novel framing
Technical Novelty 20%	Simply implements a well-known attack or defense verbatim	Applies existing technique to a new model/context	Non-trivial modification or combination of known methods	Meaningfully new approach; not obvious from prior work	Original contribution; team discovered something new
Effectiveness & Evidence 25%	Claims made without evidence or reproducible results	A few examples shown; not systematic	Consistent results across multiple test cases	Quantified success rate; benchmarked against baselines	Statistically sound; works across models/conditions
Responsible Disclosure & Ethics 5%	No consideration of misuse potential or disclosure	Aware of risks but no plan to address them	Dual-use discussed; some mitigations proposed	Clear disclosure policy; defensive framing throughout	Model of responsible red-teaming; could be a case study
Demo & Presentation 10%	Hard to follow what the project does	Shows it works but misses the compelling use case	Concrete demo, explains the value well	Live demo lands; judges want to try it themselves	Story + demo make the case unforgettably
Judge's Personal Rating 10%	Not interesting or compelling to me personally	Somewhat interesting but wouldn't think about it again	Solid idea; I can see the appeal	I'd tell a colleague about this	This genuinely excites me; I want to see it succeed

Submission

Please read before submitting:

- **One submission per team.** Only one team member should fill out this form. Duplicate submissions will be disqualified.
- **Submissions close at 6:00pm sharp.** Late submissions may not be accepted.
- **You must be present for judging.** At least one team member must be available for the in-person judging session starting at 6:30pm.
- **Have these ready before you start:** your deployed demo link, GitHub repo URL, and the names and emails of all team members.
- **Incomplete submissions may be disqualified.** Make sure every field is filled out - especially the demo link. A localhost link is not a valid demo.

Submission form

<https://docs.google.com/forms/d/e/1FAIpQLSeOBwkyvS2aaVz4KwVhviHp30wue4ADqmT6SMUls10tsWwing/viewform?usp=dialog>

Judging format

- **1. Submissions**
 - Participants submit their project here
<https://docs.google.com/forms/d/e/1FAIpQLSeOBwkyvS2aaVz4KwVhviHp30wue4ADqmT6SMUls10tsWwing/viewform?usp=dialog>
- **2. Track Round**
 - Each team presents privately in person to the judges.
 - Time limit: 3 minutes per team. 2 mins for judges to grade.
 - After all teams in a track have presented, judges select:
 - 1st place and 2nd place from each track.
 - These teams advance to the Finalist Round.
- **3. Finalist Round (Public Presentations)**
 - The advancing teams (2 per track; 6 total finalists) present again.
 - Presentations are in person and in front of the full audience.
 - Time limit: 3 minutes per team.
- **4. Final Judging & Awards**
 - The grand winner and track winners are announced.

Rules

1. Max of 4 participants per team, solo is also welcomed
2. No re-using personal projects
3. Project must be submitted before 6pm
4. Any form of cheating or rule-breaking disqualifies the team immediately
5. Organisers reserve the right to disqualify any project that violates the spirit of the event
6. Winners must be present at the awards to claim prizes
7. Prize fulfilment handled by sponsors directly - timelines may vary